

Lecture 8: Regression Analysis - Part 2

James Sears*

AFRE 891/991 FS 25

Michigan State University

*Parts of these slides are adapted from **“Data Science for Economists”** by Grant McDermott and **“Advanced Data Analytics”** by Nick Hagerty.

Table of Contents

1. Regression Output
2. IV Regression
3. Difference-in-Differences
4. Event Study and Dynamic Treatment Effects

Prologue

Empirical Analysis

We've spent the first half of our course building a **foundational knowledge** of data manipulation in R.

While cleaning, wrangling, and visualizing data can be fun, it's generally not the end goal - that's usually some form of **statistical/econometric analysis**.

Over the next weeks we're going to focus on the analysis side, using both **traditional** and **non-traditional** methods.

Empirical Analysis

This Lecture: Regression

- Standard regression: `lm()`
- Fixed effects regression: **fixest** and `feols`
- Formulas
- Choice of standard errors
- Regression Tables
- Visualizing regression output
- IV Regression
- Diff-in-Diff
 - Staggered Adoption
- Event Study and Sun-Abraham estimator

Next Lecture: Synthetic Control Methods

- Canonical Synthetic Control
- Synthetic Diff-in-Diff
 - Uniform Adoption
 - Staggered Adoption
- Partially Pooled Synthetic Control

Prologue

Packages we'll use today:

```
remotes::install_github("lrberge/fixest")
```

```
if (!require("pacman")) install.packages("pacman")  
pacman::p_load(broom, fixest, stargazer, tidyverse)  
  
options(scipen=999) # disable scientific notation
```

Prologue

As well, let's load data necessary to replicate results from [Taylor and Druckenmiller AER \(2022\)](#)¹

```
wl_df ← read_csv("data/zip_PANEL_data.csv") %>%  
  drop_na(county_fips, state_fips, claims)  
  
wl_ld_df ← readRDS("data/zip_LD_data.rds")  
  
coastal ← read_csv("data/coastal_zips.csv")  
  
did_df ← left_join(wl_ld_df, coastal, by = "zip") %>%  
  filter(coastal = FALSE)  
  
rm(wl_ld_df, coastal)
```

¹ There is a *lot* of great stuff in this paper and it's all super replicable and well-documented. I highly recommend checking it out if you're curious (only caution is that there are some *big* files in there).

Regression Output

Regression Output

Running regression is fun and all, but generally we don't just want it internal to R. Rather, we want some nicely formatted **regression output**

- **Regression Tables**
- **Regression Plots**

Regression Output

To see convenient ways to get regression output, let's run a model looking at the relationship between NFIP claims and

1. Developed area (hectares)
2. Indicators for developed area in each of the following bins:
 - [Min, 1st Quartile), [1st Quartile, Median), [Median, 3rd Quartile), [3rd Quartile, Max]
3. We'll add in state fixed effects and control for population, housing value, and income
4. Robust SE

Q: Using `feols()`, how can we write this?

Regression Output

We *could* write a bunch of code using **conditional logic**:

```
wl_df <- mutate(wl_df,
  dev_area_1 = ifelse(developed_ha <
    quantile(developed_ha, 0.25), 1,0),
  dev_area_2 = ifelse(developed_ha ≥
    quantile(developed_ha, 0.25) & developed_ha < quantile(developed_ha,
    0.5), 1,0),
  dev_area_3 = ifelse(developed_ha ≥
    quantile(developed_ha, 0.5) & developed_ha < quantile(developed_ha,
    0.75), 1,0),
  dev_area_4 = ifelse(developed_ha ≥
    quantile(developed_ha, 0.75), 1,0)
)
```

```
select(wl_df, developed_ha, starts_with("dev_")) %>%
  summarise(across(where(is.numeric), c(min, max)))
```

```
## # A tibble: 1 × 10
```

```
##   developed_ha_1 developed_ha_2 dev_area_1_1 dev_area_1_2 dev_area_2_1 1 / 104
```

```
##   <dbl> <dbl> <dbl> <dbl> <dbl>
```

Bins

Alternatively, we could use the `percent_rank()` function in **tidyverse**:

```
wl_df <- mutate(wl_df,  
                dev_area_1b = ifelse(percent_rank(developed_ha) < 0.25,  
1,0),  
                dev_area_2b = ifelse(percent_rank(developed_ha) ≥ 0.25  
& percent_rank(developed_ha) < 0.5, 1,0)  
)
```

```
identical(wl_df$dev_area_1, wl_df$dev_area_1b)
```

```
## [1] TRUE
```

```
identical(wl_df$dev_area_2, wl_df$dev_area_2b)
```

```
## [1] TRUE
```

Regression Output

1. Wetland area (hectares)
2. Indicators for developed area in each of the following bins:
 - [Min, 1st Quartile), [1st Quartile, Median), [Median, 3rd Quartile), [3rd Quartile, Max]
3. We'll add in state fixed effects and control for population, housing value, and income:
4. Robust SE

Let's specify our regression:

```
out_reg <- feols(claims ~ wetland_ha + dev_area_1 + dev_area_2 +  
                  dev_area_3 + dev_area_4 + population +  
                  housing_value + income |  
                  state_fips, data = wl_df, vcov = "HC1")
```

Regression Output

Looking at our regression output:

```
summary(out_reg)

## OLS estimation, Dep. Var.: claims
## Observations: 124,386
## Fixed-effects: state_fips: 52
## Standard-errors: Heteroskedasticity-robust
##
##              Estimate    Std. Error   t value      Pr(>|t|)
## wetland_ha      -5.664510      3.038201  -1.864429  0.062263805352 .
## dev_area_1     -76130.485716  48398.803102  -1.572983  0.115725341261
## dev_area_2     -95947.564283  46054.100961  -2.083366  0.037219880682 *
## dev_area_3    -110658.734494  43580.480628  -2.539181  0.011112435541 *
## population        7.727624      1.632395   4.733920  0.0000002204670 ***
## housing_value     0.638841      0.118570   5.387890  0.0000000071419 ***
## income          -0.005189      0.585438  -0.008864  0.992927660910
## ... 1 variable was removed because of collinearity (dev_area_4)
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## RMSE: 3,217,539.3      Adj. R2: 0.016317
```

Bins

We could also use the `bin()` function to generate a **custom binned variable**.

```
bin(x, bin)
```

- `x`: the variable/vector to divide into bins
- `bin` the list of values to split into bins

Bins

```
bin(x, bin)
```

`bin` we can specify as

- **one-sided formulas**

- yields a factor with provided labels

```
int <- 1:100  
bin(int, list("<2" = ~ x < 2, "[2,50)" = ~ x ≥ 2 & x < 50, "≥ 50" = ~ x  
> 50))
```

```
## [1] <2      [2,50) [2,50) [2,50) [2,50) [2,50) [2,50) [2,50) [2,50)  
[2,50)
```

```
## [11] [2,50) [2,50) [2,50) [2,50) [2,50) [2,50) [2,50) [2,50) [2,50)  
[2,50)
```

```
## [21] [2,50) [2,50) [2,50) [2,50) [2,50) [2,50) [2,50) [2,50) [2,50)  
[2,50)
```

```
## [31] [2,50) [2,50) [2,50) [2,50) [2,50) [2,50) [2,50) [2,50) [2,50)  
[2,50)
```

```
## [41] [2,50) [2,50) [2,50) [2,50) [2,50) [2,50) [2,50) [2,50) [2,50) 50
```


Bins

```
bin(x, bin)
```

"bin::digit" for grouping every digit interval

- Yields a numeric vector of the binned values

```
num ← seq(1, 6, 0.1)
bin(num, bin = "bin::2")
```

```
## [1] 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 4 4 4
4 4 4 4 4
## [39] 4 4 4 4 4 4 4 4 4 4 4 4 4 6
```

Bins

```
bin(x, bin)
```

`"!bin::elem"` for grouping every `elem` elements together

- Starts from **the first element**
- Yields a numeric vector of the binned values

```
bin(num, bin = "!bin::2")
```

```
## [1] 1.0 1.0 1.2 1.2 1.4 1.4 1.6 1.6 1.8 1.8 2.0 2.0 2.2 2.2 2.4 2.4  
2.6 2.6 2.8  
## [20] 2.8 3.0 3.0 3.2 3.2 3.4 3.4 3.6 3.6 3.8 3.8 4.0 4.0 4.2 4.2 4.4  
4.4 4.6 4.6  
## [39] 4.8 4.8 5.0 5.0 5.2 5.2 5.4 5.4 5.6 5.6 5.8 5.8 6.0
```

Bins

```
bin(x, bin)
```

`"!! bin::elem"` for grouping every `elem` elements together

- Starts from **the last element**
- Yields a numeric vector of the binned values

```
bin(num, bin = "!! bin::2")
```

```
## [1] 1.0 1.1 1.1 1.3 1.3 1.5 1.5 1.7 1.7 1.9 1.9 2.1 2.1 2.3 2.3 2.5  
2.5 2.7 2.7  
## [20] 2.9 2.9 3.1 3.1 3.3 3.3 3.5 3.5 3.7 3.7 3.9 3.9 4.1 4.1 4.3 4.3  
4.5 4.5 4.7  
## [39] 4.7 4.9 4.9 5.1 5.1 5.3 5.3 5.5 5.5 5.7 5.7 5.9 5.9
```

Bins

```
bin(x, bin)
```

"cut :: n" for cutting into **n** **equal parts**

- Yields a factor

```
bin(num, bin = "cut :: 4")
```

```
## [1] [1.0; 2.2] [1.0; 2.2] [1.0; 2.2] [1.0; 2.2] [1.0; 2.2] [1.0; 2.2] [1.0; 2.2]
## [7] [1.0; 2.2] [1.0; 2.2] [1.0; 2.2] [1.0; 2.2] [1.0; 2.2] [1.0; 2.2] [1.0; 2.2]
## [13] [1.0; 2.2] [2.3; 3.5] [2.3; 3.5] [2.3; 3.5] [2.3; 3.5] [2.3; 3.5] [2.3; 3.5]
## [19] [2.3; 3.5] [2.3; 3.5] [2.3; 3.5] [2.3; 3.5] [2.3; 3.5] [2.3; 3.5] [2.3; 3.5]
## [25] [2.3; 3.5] [2.3; 3.5] [3.6; 4.7] [3.6; 4.7] [3.6; 4.7] [3.6; 4.7] [3.6; 4.7]
## [31] [3.6; 4.7] [3.6; 4.7] [3.6; 4.7] [3.6; 4.7] [3.6; 4.7] [3.6; 4.7] [3.6; 4.7]
## [37] [3.6; 4.7] [3.6; 4.7] [4.8; 6.0] [4.8; 6.0] [4.8; 6.0] [4.8; 6.0] [4.8; 6.0]
## [43] [4.8; 6.0] [4.8; 6.0] [4.8; 6.0] [4.8; 6.0] [4.8; 6.0] [4.8; 6.0] [4.8; 6.0]
## [49] [4.8; 6.0] [4.8; 6.0] [4.8; 6.0]
## Levels: [1.0; 2.2] [2.3; 3.5] [3.6; 4.7] [4.8; 6.0]
```

Bins

```
bin(x, bin)
```

```
"cut :: a]b[" for cutting at values a and b
```

- Here, creates [min, a], (a-b), [b-max]
- Include as many [or] as you want/need
- Yields a factor

```
bin(num, bin = "cut::1.2]2.4[")
```

```
## [1] [1.0; 1.2] [1.0; 1.2] [1.0; 1.2] [1.3; 2.3] [1.3; 2.3] [1.3; 2.3]
## [7] [1.3; 2.3] [1.3; 2.3] [1.3; 2.3] [1.3; 2.3] [1.3; 2.3] [1.3; 2.3]
## [13] [1.3; 2.3] [1.3; 2.3] [2.4; 6.0] [2.4; 6.0] [2.4; 6.0] [2.4; 6.0]
## [19] [2.4; 6.0] [2.4; 6.0] [2.4; 6.0] [2.4; 6.0] [2.4; 6.0] [2.4; 6.0]
## [25] [2.4; 6.0] [2.4; 6.0] [2.4; 6.0] [2.4; 6.0] [2.4; 6.0] [2.4; 6.0]
## [31] [2.4; 6.0] [2.4; 6.0] [2.4; 6.0] [2.4; 6.0] [2.4; 6.0] [2.4; 6.0]
## [37] [2.4; 6.0] [2.4; 6.0] [2.4; 6.0] [2.4; 6.0] [2.4; 6.0] [2.4; 6.0]
## [43] [2.4; 6.0] [2.4; 6.0] [2.4; 6.0] [2.4; 6.0] [2.4; 6.0] [2.4; 6.0]
## [49] [2.4; 6.0] [2.4; 6.0] [2.4; 6.0]
```

Bins

```
bin(x, bin)
```

Can alternatively use `p0-p100` or `q0-q4` for cutting at percentiles/quartiles:

```
bin(num, bin = "cut::p20]p40]p60]p80]")
```

```
## [1] [1.0; 2.0] [1.0; 2.0] [1.0; 2.0] [1.0; 2.0] [1.0; 2.0] [1.0; 2.0] [1.0; 2.0]
## [7] [1.0; 2.0] [1.0; 2.0] [1.0; 2.0] [1.0; 2.0] [1.0; 2.0] [1.0; 2.0] [2.1; 3.0]
## [13] [2.1; 3.0] [2.1; 3.0] [2.1; 3.0] [2.1; 3.0] [2.1; 3.0] [2.1; 3.0] [2.1; 3.0]
## [19] [2.1; 3.0] [2.1; 3.0] [2.1; 3.0] [2.1; 3.0] [3.1; 4.0] [3.1; 4.0] [3.1; 4.0]
## [25] [3.1; 4.0] [3.1; 4.0] [3.1; 4.0] [3.1; 4.0] [3.1; 4.0] [3.1; 4.0] [3.1; 4.0]
## [31] [3.1; 4.0] [4.1; 5.0] [4.1; 5.0] [4.1; 5.0] [4.1; 5.0] [4.1; 5.0] [4.1; 5.0]
## [37] [4.1; 5.0] [4.1; 5.0] [4.1; 5.0] [4.1; 5.0] [4.1; 5.0] [4.1; 5.0] [5.1; 6.0]
## [43] [5.1; 6.0] [5.1; 6.0] [5.1; 6.0] [5.1; 6.0] [5.1; 6.0] [5.1; 6.0] [5.1; 6.0]
## [49] [5.1; 6.0] [5.1; 6.0] [5.1; 6.0] [5.1; 6.0]
## Levels: [1.0; 2.0] [2.1; 3.0] [3.1; 4.0] [4.1; 5.0] [5.1; 6.0]
```

Bins

```
bin(x, bin)
```

Can alternatively use `p0-p100` or `q0-q4` for cutting at percentiles/quartiles:

```
quantile(num, 0.5)
```

```
## 50%
```

```
## 3.5
```

```
bin(num, bin = "cut::q2[q3]")
```

```
## [1] [1.0; 3.4] [1.0; 3.4] [1.0; 3.4] [1.0; 3.4] [1.0; 3.4] [1.0; 3.4] [1.0; 3.4]
## [7] [1.0; 3.4] [1.0; 3.4] [1.0; 3.4] [1.0; 3.4] [1.0; 3.4] [1.0; 3.4] [1.0; 3.4]
## [13] [1.0; 3.4] [1.0; 3.4] [1.0; 3.4] [1.0; 3.4] [1.0; 3.4] [1.0; 3.4] [1.0; 3.4]
## [19] [1.0; 3.4] [1.0; 3.4] [1.0; 3.4] [1.0; 3.4] [1.0; 3.4] [1.0; 3.4] [1.0; 3.4]
## [25] [1.0; 3.4] [3.5; 4.7] [3.5; 4.7] [3.5; 4.7] [3.5; 4.7] [3.5; 4.7] [3.5; 4.7]
## [31] [3.5; 4.7] [3.5; 4.7] [3.5; 4.7] [3.5; 4.7] [3.5; 4.7] [3.5; 4.7] [3.5; 4.7]
## [37] [3.5; 4.7] [3.5; 4.7] [4.8; 6.0] [4.8; 6.0] [4.8; 6.0] [4.8; 6.0] [4.8; 6.0]
## [43] [4.8; 6.0] [4.8; 6.0] [4.8; 6.0] [4.8; 6.0] [4.8; 6.0] [4.8; 6.0] [4.8; 6.0]
## [49] [4.8; 6.0] [4.8; 6.0] [4.8; 6.0] [4.8; 6.0] [4.8; 6.0] [4.8; 6.0] [4.8; 6.0]
```

Bins

Using `bin()` to add a binned quartile developed area variable:

```
wl_df <- mutate(wl_df, dev_bins = bin(developed_ha,  
  "cut::p25[p50[p75[p100]]")  
  head(wl_df$dev_bins)
```

```
## [1] [0.000; 77.700] [0.000; 77.700] [0.000; 77.700] [0.000; 77.700]  
## [5] [0.000; 77.700] [0.000; 77.700]  
## 4 Levels: [0.000; 77.700] [77.702; 258.417] ... [793.156; 7339.131]
```

```
quantile(wl_df$developed_ha, 0.25)
```

```
##      25%  
## 77.70124
```


Bins

And then use the `i()` helper function in `feols()` to treat the variable as a factor:

- i.e. split the variable into a dummy for each level, minus one

```
out_reg ← feols(claims ~ wetland_ha +  
                i(dev_bins) +  
                population + housing_value + income | state_fips, data = wl_df, vcov =  
                "hetero")  
coeftable(out_reg)
```

##	Estimate	Std. Error	t value
## wetland_ha	-5.664510197	3.0382007	-1.864429234
## dev_bins::[77.702; 258.417]	-19817.078567156	5436.9887682	-3.644862885
## dev_bins::[258.421; 793.151]	-34528.248778833	15706.8991109	-2.198285514
## dev_bins::[793.156; 7339.131]	76130.485715528	48398.8031023	1.572982818
## population	7.727624318	1.6323945	4.733919570
## housing_value	0.638840970	0.1185698	5.387890316
## income	-0.005189323	0.5854383	-0.008863996
##	Pr(> t)		
## wetland_ha	0.06226380535195		
## dev_bins::[77.702; 258.417]	0.00026764332373		
## dev_bins::[258.421; 793.151]	0.02793060219016		
## dev_bins::[793.156; 7339.131]	0.11572534126324		

Bins

We can manually **specify the reference category** by adding the `ref` argument and the variable value to omit:

```
feols(claims ~ wetland_ha +  
      i(dev_bins, ref = "[793.156; 7339.131]") +  
      population + housing_value + income | state_fips, data = wl_df, vcov =  
      "hetero") %>%  
      summary()
```

```
## OLS estimation, Dep. Var.: claims  
## Observations: 124,386  
## Fixed-effects: state_fips: 52  
## Standard-errors: Heteroskedasticity-robust  
##  
##               Estimate   Std. Error   t value  
## wetland_ha          -5.664510     3.038201 -1.864429  
## dev_bins::[0.000; 77.700]    -76130.485716 48398.803102 -1.572983  
## dev_bins::[77.702; 258.417]  -95947.564283 46054.100961 -2.083366  
## dev_bins::[258.421; 793.151] -110658.734494 43580.480628 -2.539181  
## population           7.727624     1.632395  4.733920  
## housing_value         0.638841     0.118570  5.387890  
## income              -0.005189     0.585438 -0.008864  
##  
##               Pr(>|t|)  
## wetland_ha          0.062263805352 .  
## dev_bins::[0.000; 77.700]    0.115725341261  
## dev_bins::[77.702; 258.417]  0.037310000000
```

Regression Table

To format our data as a regression table we have two choices

1. Use **stargazer** on the summary dataframe (or `lm` object)
 - Works well for summary statistics tables too
2. Use `etable()` to create a well-formatted regression table
 - Works on one or multiple `fixest` objects

Stargazer

The **stargazer** package's `stargazer()` function has great functionality for obtaining tables in HTML, LaTeX, or plain text format

- Also the preferred method for regression tables with `lm()` or `feIm()` objects
- Doesn't play nice with `fixest` objects

For example, let's say you want to produce a **summary statistics table** for our included variables:

```
sumstats <- select(wl_df, claims, wetland_ha, starts_with("dev_area_")) %>%  
  stargazer(type = "text")
```

```
##  
## =====  
## Statistic N Mean St. Dev. Min Max  
## =====
```

Stargazer

This empty table reveals one restriction of `stargazer`: it requires the input be formatted as a **data frame**.

```
sumstats <- select(wl_df, claims, wetland_ha, starts_with("dev_area_"))
%>%
  as.data.frame() %>%
  stargazer(type = "text")
```

```
##
## =====
## Statistic      N      Mean      St. Dev.      Min      Max
## -----
## claims        131,588 109,330.100 3,155,095.000 0.000 418,203,035.000
## wetland_ha    131,588  1,301.656   4,533.807   0.000   261,518.300
## dev_area_1    131,588    0.250     0.433     0       1
## dev_area_2    131,588    0.250     0.433     0       1
## dev_area_3    131,588    0.250     0.433     0       1
## dev_area_4    131,588    0.250     0.433     0       1
## dev_area_1b   131,588    0.250     0.433     0       1
## dev_area_2b   131,588    0.250     0.433     0       1
```

Stargazer

We can easily get LaTeX output:

```
sumstats <- select(wl_df, claims, wetland_ha, starts_with("dev_area_"))
%>%
  as.data.frame() %>%
  stargazer(type = "latex")
```

```
##
## % Table created by stargazer v.5.2.3 by Marek Hlavac, Social Policy
## Institute. E-mail: marek.hlavac at gmail.com
## % Date and time: Mon, Nov 03, 2025 - 4:14:45 PM
## \begin{table}[!htbp] \centering
##   \caption{}
##   \label{}
## \begin{tabular}{@{\extracolsep{5pt}}lcccc}
## \\\[-1.8ex]\hline
## \hline \\\[-1.8ex]
## Statistic & \multicolumn{1}{c}{N} & \multicolumn{1}{c}{Mean} & &
## \multicolumn{1}{c}{St. Dev.} & \multicolumn{1}{c}{Min} & \multicolumn{1}{c}{Max} \\\
```

Stargazer

Or HTML:

```
sumstats <- select(wl_df, claims, wetland_ha, starts_with("dev_area_"))
%>%
  as.data.frame() %>%
  stargazer(type = "html")
```

Statistic	N	Mean	St. Dev.	Min	Max
claims	131,588	109,330.100	3,155,095.000	0.000	418,203,035.000
wetland_ha	131,588	1,301.656	4,533.807	0.000	261,518.300
dev_area_1	131,588	0.250	0.433	0	1
dev_area_2	131,588	0.250	0.433	0	1
dev_area_3	131,588	0.250	0.433	0	1
dev_area_4	131,588	0.250	0.433	0	1

Stargazer

Built-in styles also allow you to easily replicate the appearance of tables in several popular journals

```
sumstats <- select(wl_df, claims, wetland_ha, starts_with("dev_area_"))
%>%
  as.data.frame() %>%
  stargazer(style = "aer", type = "html")
```

Statistic	N	Mean	St. Dev.	Min	Max
claims	131,588	109,330.100	3,155,095.000	0.000	418,203,035.000
wetland_ha	131,588	1,301.656	4,533.807	0.000	261,518.300
dev_area_1	131,588	0.250	0.433	0	1
dev_area_2	131,588	0.250	0.433	0	1
dev_area_3	131,588	0.250	0.433	0	1
dev_area_4	131,588	0.250	0.433	0	1

Stargazer

Have a custom table you want to print out? No problem! Just use the `summary= FALSE` argument.

```
group_by(wl_df, year) %>%  
  summarise("Total Claims" = sum(claims, na.rm = T),  
            "Total Wetland Change" = sum(wetland_change, na.rm = T)) %>%  
  stargazer(summary = F, type = "html", digits = 4)
```

	year	Total Claims	Total Wetland Change
1	2001	632824948.543333	0
2	2006	6082244449.66	173394.261882353
3	2011	2361949781.01333	-492149.872235294
4	2016	5309509545.73667	311423.318117647

Stargazer

We can also use stargazer to build regression tables directly with `lm` output:

```
lm1 <- lm(claims ~ wetland_ha, data = wl_df)
lm2 <- lm(claims ~ wetland_ha + population, data = wl_df)
lm3 <- lm(claims ~ wetland_ha + population + housing_value, data = wl_df)

stargazer(lm1, lm2, lm3, style = "aer", type = "html")
```

	claims		
	(1)	(2)	(3)
wetland_ha	5.948 ^{***} (1.918)	6.758 ^{***} (1.932)	7.492 ^{***} (2.035)
population		10.203 ^{***} (0.633)	9.821 ^{***} (0.676)
housing_value			0.136 ^{**}

Stargazer

`stargazer` has a full suite of options for adding titles, notes, changing variable names, column labeling, etc.

```
stargazer(lm1, lm2, lm3, style = "aer", type = "html",
          title = "NFIP Claims and Wetland Area", # chart title
          dep.var.labels = c("NFIP Claims Amount"), # label dep.
          variable
          covariate.labels = c("Wetland Area (Ha)", "Population",
                                "Housing Value"), # label variables
          ci = c(T,T,T), # replace std errors with conf. intervals
          digits = 4, # include 4 digits
          initial.zero = T, # add initial zero for values 0-1
          label = "table_ols", # latex label
          note = c("This Table reports regression results of NFIP claims
                    amounts on wetland acreage in a given zipcode, controlling for various
                    factors."), # custom note
          out = "output/stargazer_table.html" # path to save out the
          table
          )
```

NFIP Claims and Wetland Area			
	NFIP Claims Amount		
	(1)	(2)	(3)
Wetland Area (Ha)	5.9476 ^{***}	6.7584 ^{***}	7.4920 ^{***}
	(2.1877, 9.7075)	(2.9726, 10.5442)	(3.5032, 11.4808)
Population		10.2027 ^{***}	9.8212 ^{***}
		(8.9614, 11.4439)	(8.4968, 11.1455)
Housing Value			0.1361 ^{**}
			(0.0087, 0.2636)
Constant	101,588.3000 ^{***}	4,782.7480	-14,013.8800
	(83,853.0500, 119,323.6000)	(-16,770.8700, 26,336.3700)	(-42,598.7600, 14,570.9900)

etable

While we can't use stargazer with `fixest` objects, the package provides the equivalent `etable()` function that allows for fully professional regression tables.

Suppose we wanted to show the value of our wetland coefficient across several different fixed effects specifications:

```
reg_state <- feols(claims ~ wetland_ha + i(dev_bins, ref = "[793.156;  
7339.131]") +  
                  population + housing_value + income | state_fips, data =  
wl_df, vcov = "hetero")  
reg_cty <- feols(claims ~ wetland_ha + i(dev_bins, ref = "[793.156; 7339.131]")  
+  
                  population + housing_value + income | county_fips, data =  
wl_df, vcov = "hetero")  
reg_styr <- feols(claims ~ wetland_ha + i(dev_bins, ref = "[793.156; 7339.131]")  
+  
                  population + housing_value + income | state_fips + year, data  
= wl_df, vcov = "hetero")  
reg_ctyr <- feols(claims ~ wetland_ha + i(dev_bins, ref = "[793.156; 7339.131]")  
+  
                  population + housing_value + income | state_fips + year, data  
= wl_df, vcov = "hetero")
```

etable

Putting together the regression table:

```
etable(reg_state, reg_cty, reg_styr, reg_ctyr,  
  title = "Fixed Effects Progression",  
  fixef_sizes = T,  
  digits = 4, # digits for coefficients and std. errors  
  digits.stats = 4, # digits for fit stats  
  tex = F # whether output is LaTeX (T) or a df (F)  
)
```

```
##                                reg_state                reg_cty  
## Dependent Var.:                claims                claims  
##  
## wetland_ha                    -5.665. (3.038)           -3.251 (2.529)  
## dev_bins = [0.000;77.700]      -76,130.5 (48,398.8)     -63,319.6 (57,328.2)  
## dev_bins = [77.702;258.417]    -95,947.6* (46,054.1)     -74,749.1 (53,209.3)  
## dev_bins = [258.421;793.151]  -110,658.7* (43,580.5)    -98,888.3. (53,014.5)  
## population                     7.728** (1.632)           6.101** (2.189)  
## housing_value                  0.6388** (0.1186)          0.3384** (0.1082)  
## income                        -0.0052 (0.5854)           1.676** (0.5679)
```

```
## Fixed-Effects: -----
```

```
## state_fips                    Yes                        No
```

```
## income_fips                    No                        Yes
```

etable

Another convenient feature of `etable`: we can **change standard error estimators** on the fly!

```
etable(reg_state, reg_cty, reg_styr, reg_ctyr,  
       title = "Fixed Effects Progression",  
       fixef_sizes = T,  
       digits = 4, # digits for coefficients and std. errors  
       digits.stats = 4, # digits for fit stats  
       tex = F, # whether output is LaTeX (T) or a df (F)  
       se = "standard"  
)
```

```
##                               reg_state                               reg_cty  
## Dependent Var.:                claims                               claims  
##  
## wetland_ha                    -5.665** (2.178)                    -3.251 (2.758)  
## dev_bins = [0.000;77.700]      -76,130.5* (38,511.7)      -63,319.6 (42,225.7)  
## dev_bins = [77.702;258.417]    -95,947.6** (35,421.5)    -74,749.1. (38,261.3)  
## dev_bins = [258.421;793.151] -110,658.7*** (31,613.6) -98,888.3** (33,363.8)  
## population                     7.728*** (1.017)           6.101*** (1.180)  
## housing_value                  0.6388*** (0.1042)          0.3384** (0.1215)  
## income                        -0.0052 (0.6077)             1.676* (0.6562)
```

etable

Another convenient feature of `etable`: we can **change standard error estimators** on the fly!

```
etable(reg_state, reg_cty, reg_styr, reg_ctyr,  
       title = "Fixed Effects Progression",  
       fixef_sizes = T,  
       digits = 4, # digits for coefficients and std. errors  
       digits.stats = 4, # digits for fit stats  
       tex = F, # whether output is LaTeX (T) or a df (F)  
       cluster = "state_fips"  
)
```

```
##                               reg_state                reg_cty  
## Dependent Var.:                claims                claims  
##  
## wetland_ha                    -5.665 (9.362)            -3.251 (3.040)  
## dev_bins = [0.000;77.700]      -76,130.5 (96,632.1) -63,319.6 (122,943.7)  
## dev_bins = [77.702;258.417]    -95,947.6 (97,532.5) -74,749.1 (119,884.7)  
## dev_bins = [258.421;793.151] -110,658.7 (81,641.5) -98,888.3 (104,571.2)  
## population                     7.728 (4.962)            6.101 (5.319)  
## housing_value                  0.6388. (0.3581)          0.3384 (0.2336)  
## income                        -0.0052 (1.440)           1.676 (1.210)
```


feols() and etable

More efficient: take advantage of **stepwise functions** for **multiple estimation**:

```
regs_fe ← feols(claims ~ wetland_ha + i(dev_bins, ref = "[793.156; 7339.131]")
+
  population + housing_value + income |
  sw0(state_fips, county_fips, state_fips + year, county_fips + year),
  data = wl_df, vcov = "hetero")
etable(regs_fe)
```

##	regs_fe.1	regs_fe.2
## Dependent Var.:	claims	claims
##		
## Constant	142,577.5* (65,774.1)	
## wetland_ha	6.198*** (1.792)	-5.665. (3.038)
## dev_bins = [0.000;77.700]	-181,068.7*** (53,079.4)	-76,130.5 (48,399.0)
## dev_bins = [77.702;258.417]	-183,517.4*** (50,834.8)	-95,947.6* (46,054.3)
## dev_bins = [258.421;793.151]	-165,192.3*** (46,304.8)	-110,658.7* (43,580.7)
## population	5.807*** (1.566)	7.728*** (1.632)
## housing_value	0.1360 (0.0894)	0.6388*** (0.1186)
## income	0.2983 (0.5802)	-0.0052 (0.5854)
## Fixed-Effects:	-----	-----

feols and stepwise estimation

With `feols()` we can combine **stepwise functions** on the RHS with **multiple dependent variables** on the LHS, all within the same object!

Consider the (not valid) regression of

- Claims, housing values and income on population
- Stepwise through 5 different FE specifications

```
regs_both ← feols(c(claims, housing_value, income) ~ population |  
  sw0(state_fips, county_fips, state_fips + year, county_fips + year),  
  data = wl_df, vcov = "hetero")
```

fixest_multi Objects

Which yields a `fixest_multi` object with **15 different model estimates!**

We can access its elements with five keys:

- `sample`
- `lhs`
- `rhs`
- `fixef`
- `iv`

fixest_multi Objects

We can perform **basic access** of elements using the `object[key = value]` syntax

For example, we can access the first LHS with

```
regs_both[lhs = 1]

## Standard-errors: Heteroskedasticity-robust
## Dep. var.: claims
## Fixed-effects: 1
##               Estimate Std. Error  t value  Pr(>|t|)
## (Intercept) 14233.1923 6896.03764   2.06397   0.039023 *
## population    10.1418    1.00254 10.11612 < 2.2e-16 ***
## ---
## Fixed-effects: state_fips
##               Estimate Std. Error t value  Pr(>|t|)
## population  10.7946    1.10676 9.75334 < 2.2e-16 ***
## ---
## Fixed-effects: county_fips
##               Estimate Std. Error t value  Pr(>|t|)
```

fixest_multi Objects

In this case, all our models are estimated on the same sample, and R tells us if we act a fool and behave otherwise

```
regs_both[sample = 2]
```

```
## Error: in set_index_multi(sample, index_names):  
## The index sample is not valid for this list of results (the valid ones  
are  
## fixef and lhs).  
## You can, however, set it to 1 if necessary, as in sample = 1.
```

etable and fixest_multi

We can combine this syntax with `etable()` to plot specific estimates at once:

```
etable(regs_both[fixef = 1])

##                regs_both[fixef = 1]  regs_both[fixef = 2]
regs_both[fixef = 3]
## Dependent Var.:                claims                housing_value
income
##
## Constant                14,233.2* (6,896.0) 133,103.8*** (447.7) 47,786.9***
(74.13)
## population                10.14*** (1.003)    2.807*** (0.0345) 0.2762***
(0.0048)
## -----
-----
## S.E. type                Heteroskedast.-rob. Heteroskedast.-rob.
Heteroskedast.-rob.
## Observations                129,966                125,059
126,013
```

etable and fixest_multi

The order in which we put our keys will **re-order the columns**

- default is `sample > lhs > fixef > rhs > iv`
- set key to logical to include/omit all

```
etable(regs_both[lhs = 2, fixef = TRUE])
```

```
##                regs_both[lhs = 2..1 regs_both[lhs ..2 regs_both[lhs =  
..3  
## Dependent Var.:          housing_value      housing_value  
housing_value  
##  
## Constant          133,103.8*** (447.7)  
## population          2.807*** (0.0345) 1.614*** (0.0324) -0.5511***  
(0.0388)  
## Fixed-Effects:  -----  
---  
## state_fips                No                Yes  
No  
## county_fips              No                No
```

etable and fixest_multi

The order in which we put our keys will **re-order the columns**

- default is `sample > lhs > fixef > rhs > iv`
- set key to logical to include/omit all

```
etable(regs_both[fixef = 2:4, lhs = 2])
```

```
##                regs_both[fixe..1 regs_both[fixef ..2 regs_both[fixe..3
## Dependent Var.:      housing_value      housing_value      housing_value
##
## population          1.614*** (0.0324) -0.5511*** (0.0388) 1.522*** (0.0308)
## Fixed-Effects: -----
## state_fips                Yes                No                Yes
## county_fips               No                Yes                No
## year                      No                No                Yes
## -----
## S.E. type      Heteroskedast--rob. Heteroskedast--rob. Heteroskedast--rob.
## Observations           125,059           125,059           125,059
## R2                   0.33377           0.59430           0.38778
## Within R2            0.03016           0.00364           0.02916
```


etable

The `dict` argument also makes it straightforward to **add well-formatted variable names**:

```
etable(regs_fe,
      dict = c(wetland_ha = "Wetland Area (Ha)",
                dev_bins = "Developed Area",
                state_fips = "State",
                county_fips = "County",
                year = "Year"))
```

```
##                                                    regs_fe.1
## Dependent Var.:                                claims
##
## Constant                                     142,577.5* (65,774.1)
## Wetland Area (Ha)                           6.198*** (1.792)
## Developed Area = [0.000;77.700]             -181,068.7*** (53,079.4)
## Developed Area = [77.702;258.417]           -183,517.4*** (50,834.8)
## Developed Area = [258.421;793.151]          -165,192.3*** (46,304.8)
## population                                   5.807*** (1.566)
## housing_value                               0.1360 (0.0894)
## income                                       0.2983 (0.5802)
## Fixed-Effects: -----
```

etable

You can also set up a **default, external dictionary** with `setFixest_dict()` that will automatically be used in all following tables.

```
# set dictionary for use across the entire script
setFixest_dict(c(wetland_ha = "Wetland Area (Ha)",
                 dev_bins = "Developed Area",
                 state_fips = "State",
                 county_fips = "County",
                 year = "Year"))
```

etable

You can also set up a **default, external dictionary** with `setFixest_dict()` that will automatically be used in all following tables.

```
# Variables are automatically re-labeled
etable(regs_fe)

##                                                    regs_fe.1
## Dependent Var.:                                claims
##
## Constant                                     142,577.5* (65,774.1)
## Wetland Area (Ha)                           6.198*** (1.792)
## Developed Area = [0.000;77.700]             -181,068.7*** (53,079.4)
## Developed Area = [77.702;258.417]           -183,517.4*** (50,834.8)
## Developed Area = [258.421;793.151]          -165,192.3*** (46,304.8)
## population                                   5.807*** (1.566)
## housing_value                               0.1360 (0.0894)
## income                                       0.2983 (0.5802)
## Fixed-Effects:  -----
## State                                           No
## County                                           No
```

This is only a **fraction of the settings** available. Don't worry, we'll get more practice with them in the remainder of the lecture (and in the next problem set)!

IV Regression

IV Regression

`fixest` also makes it easy to estimate IV regressions.

All we need is to add our IV formula after a second `|`

```
feols(y ~ x_1 + ... + X_n | fe_1 + ... + fe_n
```

```
| x_endo1 + x_endo2 ~ x_inst1 + x_inst2, data = df)
```

- `x_endo`: the endogenous regressor(s)
- `x_inst`: the instrument(s)

IV Regression

Let's suppose that wetland area in 2016 was endogenous due to recent policies, but that the area in 2001 was a valid instrument³. Using the `did_df` data:

```
iv_1 <- feols(claims_2016_3yr ~ developed_ha_2016 + population_2016 +  
              income_2016 + housing_value_2016 | # dep var ~ expl vars  
              state_fips | # FE  
              wetland_ha_2016 ~ wetland_ha_2001, #<< # endog. var ~  
              instrument  
              data = did_df)
```

³ This is purely a toy example - there isn't an endogeneity issue in the paper.

IV Regression

By default, `summary()` now reports the **second stage results**.

```
summary(iv_1)

## TSLS estimation - Dep. Var.: claims_2016_3yr
##                               Endo.      : wetland_ha_2016
##                               Instr.     : wetland_ha_2001
## Second stage: Dep. Var.: claims_2016_3yr
## Observations: 25,006
## Fixed-effects: state_fips: 49
## Standard-errors: Clustered (state_fips)
##                               Estimate Std. Error  t value Pr(>|t|)
## fit_wetland_ha_2016      9.024196    8.350857   1.08063  0.28526
## developed_ha_2016    103.424292   90.923739   1.13748  0.26098
## population_2016      14.942096   10.430715   1.43251  0.15848
## income_2016           4.731703    3.268778   1.44754  0.15424
## housing_value_2016    -0.402308    0.273227  -1.47243  0.14743
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## RMSE: 2,652,351.7      Adj. R2: 0.04286
```


IV Regression

We can get **first-stage results** with the `stage = 1` argument:

```
summary(iv_1, stage = 1)

## TSLS estimation - Dep. Var.: wetland_ha_2016
##               Endo.       : wetland_ha_2016
##               Instr.      : wetland_ha_2001
## First stage: Dep. Var.: wetland_ha_2016
## Observations: 25,006
## Fixed-effects: state_fips: 49
## Standard-errors: Clustered (state_fips)
##               Estimate Std. Error   t value  Pr(>|t|)
## wetland_ha_2001   1.000634   0.000810 1235.28173 < 2.2e-16 ***
## developed_ha_2016 -0.012556   0.004855  -2.58637  0.012788 *
## population_2016   0.000255   0.000206   1.23581  0.222544
## income_2016       -0.000074   0.000039  -1.86483  0.068325 .
## housing_value_2016 -0.000015   0.000011  -1.43371  0.158139
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## RMSE: 102.7      Adj. R2: 0.999458
```

IV Regression

Or both at the same time with `stage = 2:1` or `1:2`

```
summary(iv_1, stage = 1:2)
```

```
## IV: First stage: wetland_ha_2016
## TSLS estimation - Dep. Var.: wetland_ha_2016
##               Endo.      : wetland_ha_2016
##               Instr.     : wetland_ha_2001
## First stage: Dep. Var.: wetland_ha_2016
## Observations: 25,006
## Fixed-effects: state_fips: 49
## Standard-errors: Clustered (state_fips)
##               Estimate Std. Error   t value  Pr(>|t|)
## wetland_ha_2001    1.000634   0.000810 1235.28173 < 2.2e-16 ***
## developed_ha_2016  -0.012556   0.004855  -2.58637  0.012788 *
## population_2016     0.000255   0.000206   1.23581  0.222544
## income_2016        -0.000074   0.000039  -1.86483  0.068325 .
## housing_value_2016 -0.000015   0.000011  -1.43371  0.158139
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## RMSE: 102.7      Adj. R2: 0.999458
##               Within R2: 0.999375
## F-test (1st stage): stat = 38,825,154.0, p < 2.2e-16, on 1 and 25,000 DoF.
##
## IV: Second stage
```

IV Regression

And the output table (adding a few IV fit statistics with `fitstat`):

```
etable(iv_1, stage = 1:2, fitstat = ~ . + ivfall + ivwaldall)
```

```
##                               iv_1.1          iv_1.2
## IV stages                     First          Second
## Dependent Var.:      wetland_ha_2016  claims_2016_3yr
##
## wetland_ha_2001      1.001*** (0.0008)
## developed_ha_2016    -0.0126* (0.0049)    103.4 (90.92)
## population_2016      0.0003 (0.0002)     14.94 (10.43)
## income_2016          -7.36e-5. (3.95e-5)   4.732 (3.269)
## housing_value_2016   -1.52e-5 (1.06e-5) -0.4023 (0.2732)
## wetland_ha_2016                                9.024 (8.351)
## Fixed-Effects:      -----
## State                                Yes          Yes
## -----
## S.E.: Clustered      by: State          by: State
## Observations          25,006            25,006
## R2                    0.99946            0.04489
## Within R2             0.99938            0.01183
## F-test (IV only)      38,825,154.0        4.7368
## Wald (IV only)        1,525,921.0         1.1678
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Difference-in-Differences

Difference-in-Differences

Let's now replicate some actual results from the paper: the **upstream/downstream diff-in-diff**.

$$\Delta F_{is} = \beta \Delta W_{is} + \gamma \Delta W_{is}^{UP} + \lambda \Delta W_{is}^{ALL} + \theta \Delta X_{is} + \alpha_s + \epsilon_{is}$$

- ΔF_{is} : the change in flood insurance claims from 2001 to 2016
- ΔW_{is} : the change in **within-zipcode** wetland area from 2001 to 2016
- ΔW_{is}^{UP} : the change in **upstream** wetland area from 2001 to 2016
- ΔW_{is}^{ALL} : the change in **upstream and downstream** wetland area from 2001 to 2016
- ΔX_{is} : the change in time-varying covariates from 2001 to 2016
- α_s state fixed effects

Difference-in-Differences

$$\Delta F_{is} = \beta \Delta W_{is} + \gamma \Delta W_{is}^{UP} + \lambda \Delta W_{is}^{ALL} + \theta \Delta X_{is} + \alpha_s + \epsilon_{is}$$

Diff-in-diff style coefficients of interest are

- β : the effect of "local" wetlands (within the same zipcode)
- γ : the differential effect of upstream wetlands
 - i.e. the "direct protective services" of the wetland

Difference-in-Differences

$$\Delta F_{is} = \beta \Delta W_{is} + \gamma \Delta W_{is}^{UP} + \lambda \Delta W_{is}^{ALL} + \theta \Delta X_{is} + \alpha_s + \epsilon_{is}$$

```
did_reg <- feols(claims_2001_2016 ~ wetland_2001_2016 +  
  wetlands_change_up + wetlands_change_all +  
  developed_2001_2016 +  
  income_2001_2016 +  
  population_2001_2016 +  
  housing_units_2001_2016 +  
  housing_value_2001_2016 +  
  CRS_2001_2016  
  |state_fips, cluster = ~county_fips, data = did_df)
```

Difference-in-Differences

$$\Delta F_{is} = \beta \Delta W_{is} + \gamma \Delta W_{is}^{UP} + \lambda \Delta W_{is}^{ALL} + \theta \Delta X_{is} + \alpha_s + \epsilon_{is}$$

```
coefTable(did_reg)
```

##	Estimate	Std. Error	t value
Pr(> t)			
## wetland_2001_2016 0.122423316	-157.7758561	102.112020	-1.5451252
## wetlands_change_up 0.018290952	-499.9636949	211.760190	-2.3609900
## wetlands_change_all 0.168607941	-21.0437746	15.282017	-1.3770287
## developed_2001_2016 0.042948914	337.5158492	166.667757	2.0250818
## income_2001_2016 0.638967340	-0.5575620	1.188346	-0.4691917
## population_2001_2016 0.624631357	-5.5928337	11.429141	-0.4893486
## housing_units_2001_2016 0.008653789	76.3330446	29.054958	2.6271951

Difference-in-Differences

Comparing to the published paper results (Column 2):

TABLE 1—THE EFFECT OF WETLANDS ON FLOOD DAMAGES						
	Zip code-level NFIP claims (US\$)					
	LD (1)	DID (2)	Panel (3)	LD (4)	DID (5)	Panel (6)
<i>Wetland effects</i>						
Local wetland change (hectares, or ha)	−229.2 (127.7)	−157.8 (102.1)	−180.9 (83.6)			
Local wetland gain (ha)				−24.1 (116.4)	39.0 (74.7)	153.6 (220.9)
Local wetland loss (ha)				−495.3 (250.8)	−450.8 (247.2)	−461.7 (272.4)
Upstream wetland change (ha)		−500.0 (211.8)				
Upstream wetland gain (ha)					−71.9 (77.6)	
Upstream wetland loss (ha)					−810.4 (342.0)	
<i>Controls</i>						
Developed area (ha)	390.7 (172.0)	337.5 (166.7)	2,863.7 (2,167.9)	372.0 (170.1)	312.0 (165.7)	2,866.5 (2,170.0)
Median income (US\$)	−0.5 (1.1)	−0.6 (1.2)	1.0 (2.2)	−0.5 (1.1)	−0.5 (1.2)	1.0 (2.2)
Population	−6.9 (12.1)	−5.6 (11.4)	−165.6 (167.1)	−6.4 (12.0)	−5.0 (11.4)	−165.3 (167.0)
Number of housing units	77.5 (25.5)	76.3 (29.1)	339.1 (204.5)	76.8 (25.4)	75.7 (28.9)	337.8 (204.5)
Median home value (US\$)	0.2 (0.2)	0.2 (0.3)	0.6 (0.5)	0.3 (0.2)	0.2 (0.3)	0.6 (0.5)
CRS discount (percent)	159,307 (100,644)	169,618 (117,832)	135,364 (111,952)	159,242 (100,657)	168,208 (117,592)	135,424 (111,940)

Difference-in-Differences

Historically, we've thought mostly about whether **parallel trends** are likely to hold for validity of our difference-in-differences estimator.

Logic:

- Control and treated groups are trending similarly prior to treated group receiving treatment
- Ensures validity of control group's post-treatment trend as counterfactual for treatment group absent the treatment

Most often (and especially with staggered adoption timing), we estimate these models in the canonical **two-way fixed effects** specification (**TWFE**)

$$Y_{it} = \beta T_{it} + \delta_t + \alpha_i + \epsilon_{it}$$

Difference-in-Differences

However, recent literature has shown that even if **parallel trends** hold, our estimates might **still be biased**

- Due to unintended weighting of all possible 2x2 comparisons (**Goodman-Bacon (2021)**)
- Which may even be negative (**de Chaisemartin and D'Haultfoeuille (2020)**)

We'll get into some **alternate estimators** that overcome these limitations soon (and next lecture), but for now let's chat briefly about another situation that warrants a different empirical specification: **dynamic treatment effects**

Event Study and Dynamic Treatment Effects

Dynamic Treatment Effects

Recall the TWFE diff-in-diff model:

$$Y_{it} = \beta T_{it} + \delta_t + \alpha_i + \epsilon_{it}$$

This yields an estimate of β , the **average treatment effect for the treated (ATT)**.

- Average across all treated units
- Average across all post-treatment time periods

However, there are plenty of settings where we might expect **dynamic treatment effects**: treatment effects that differ by the amount of time since receiving treatment.

Event Study

This is where **event study**⁴ comes in.

Idea: instead of a single $\hat{\beta}$, estimate a vector of treatment effects $\hat{\beta}^k$

- Each coefficient estimates a period-specific ATT for being k time periods relative to receiving treatment
 - i.e. ATT 1 day post-treatment may look quite different than 100 days post-treatment

Estimating this requires thinking about time in an alternate fashion (no, this isn't a multiverse situation)...

⁴ Note that the "original" event study methods refer to techniques in finance related to the impact of various "events" on stock prices, generally for a single asset. The event studies we're talking about today refer to the approach used in applied economics for estimating dynamic treatment paths.

Time-to-Event Data

To work through this, let's use some data I'm pretty familiar with: COVID-19 stay-at-home mandate data.

```
sah ← read_csv("data/sah.csv") %>%  
  drop_na(cadt) %>%  
  select(state, date, weekday, cadt, visits, mandate_date)
```

```
head(sah)
```

```
## # A tibble: 6 × 6  
##   state date      weekday  cadt visits mandate_date  
##   <chr> <date>      <dbl> <dbl> <dbl> <date>  
## 1 AK    2020-02-24      2  3.19  -2.14 2020-03-29  
## 2 AK    2020-02-25      3  2.38   2.01 2020-03-29  
## 3 AK    2020-02-26      4  7.36   2.94 2020-03-29  
## 4 AK    2020-02-27      5  6.97   5.25 2020-03-29  
## 5 AK    2020-02-28      6  6.40   6.50 2020-03-29  
## 6 AK    2020-02-29      7  1.53  10.0 2020-03-29
```

Aside: Diff-in-Diff

Here we have information on the unit (`state`), the time period (`date`), and the date of treatment (`mandate_date`).

- States adopted mandates at different times $\Rightarrow$ **staggered adoption**

To estimate the TWFE diff-in-diff, we just need a treatment indicator

- = 1 for treated units after treatment occurred
 - mandate-adopting states post-mandate adoption
- = 0 for treated units prior to treatment
 - mandate-adopting states prior to mandate adoption
- = 0 for control units
 - states that never adopted a statewide stay-at-home mandate

Aside: Diff-in-Diff

To estimate the TWFE diff-in-diff, we just need a treatment indicator.

First, identify which states ever adopted a stay-at-home mandate:

```
sah_did <- group_by(sah, state) %>%  
  # get state-specific max value of mandate adoption date  
  dplyr::mutate(sah_max = max(mandate_date, na.rm = T))
```

This should yield a date for adopting states, and `-Inf` for non-adopters (i.e. pure controls)

```
unique(sah_did$sah_max)
```

```
## [1] "2020-03-29" "2020-04-05" "-Inf"      "2020-04-01" "2020-03-19"  
## [6] "2020-03-26" "2020-03-24" "2020-04-03" "2020-03-25" "2020-03-22"  
## [11] "2020-03-30" "2020-03-27" "2020-03-31" "2020-04-02" "2020-03-28"  
## [16] "2020-04-06" "2020-04-04" "2020-03-23" "2020-04-08"
```

Aside: Diff-in-Diff

Next, determine whether we are in the post-adoption period for our adopting states (i.e. our treatment dummy):

```
# note: data are still grouped by state
sah_did <- dplyr::mutate(sah_did,
  post_treat = case_when(
    sah_max == -Inf ~ 0, # 0 for pure controls
    date ≥ mandate_date ~ 1, # 1 for adopting states in post-period
    TRUE ~ 0 # 0 for adopting states pre-treatment
  )
) %>%
ungroup()

summary(sah_did$post_treat)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##  0.00000  0.00000  0.00000  0.4204  1.00000  1.00000
```

Aside: Diff-in-Diff

Estimating the TWFE diff-in-diff for the impact of stay-at-home on average distance traveled:

```
did_1 <- feols(cadt ~ post_treat | state + date, data = sah_did, cluster
= ~ state)
summary(did_1)
```

```
## OLS estimation, Dep. Var.: cadt
## Observations: 3,366
## Fixed-effects: state: 51,  date: 66
## Standard-errors: Clustered (state)
##               Estimate Std. Error  t value      Pr(>|t|)
## post_treat -5.50772      1.04595 -5.26575 0.0000029507 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## RMSE: 5.22448      Adj. R2: 0.929189
##               Within R2: 0.059987
```

Aside: Diff-in-Diff

Comparing to results from [Sears et al. \(2023\)](#) (Column 5):

TABLE 2. Statewide stay-at-home mandates, travel activity, and social distancing

	Early SAH states				All SAH states			
	<i>ΔDT</i> (1)	<i>ΔDT</i> (2)	<i>NEV</i> (3)	<i>ENC</i> (4)	<i>ΔDT</i> (5)	<i>ΔDT</i> (6)	<i>NEV</i> (7)	<i>ENC</i> (8)
SAH_{it}	-4.454 ^b (2.074)	2.495 (2.113)	2.629 (1.927)	3.050 (4.598)	-5.508 ^a (1.036)	-6.992 ^a (1.454)	-2.149 ^b (0.880)	-3.506 ^a (0.981)
$\bar{Y}$	-26.59	-29.29	-36.23	-53.57	-26.59	-29.29	-36.23	-53.57
Pre-period $\bar{Y}$	-3.98	-4.98	-7.02	-14.85	-3.98	-4.98	-7.02	-14.85
State + date FE	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes
Cohort pre-trends	No	Yes	Yes	Yes	No	Yes	Yes	Yes
N	3,366	3,366	3,366	3,300	3,366	3,366	3,366	3,300
Adjusted R^2	0.926	0.896	0.955	0.954	0.929	0.901	0.955	0.954

Note: Standard errors are clustered at the state level. ^a $p < 0.01$, ^b $p < 0.05$.

Diff-in-Diff: Weighting Issues

If we decompose the overall ATT into each 2x2 estimator, we can get more information as to which comparisons are driving our treatment effects:

```
pacman::p_load(bacondecomp)

# first, get outcome residualized of fixed effects
sah_did <- sah_did %>%
  dplyr::mutate(cadt_r = feols(cadt ~ 0 | state + date, data =
sah_did)$residuals,
               date_num = as.numeric(date))

# run decomp of residualized outcome on treatment
decomp <- bacon(cadt_r ~ post_treat,
               data = sah_did,
               id_var = "state",
               time_var = "date_num",
               quietly = T) %>%
  mutate(full_att = sum(estimate*weight))
```

Aside: Bacon Decomposition

Looking at the different weights:

```
decomp_tab ← decomp %>%
  dplyr::select(type, weight, estimate, full_att) %>%
  dplyr::mutate(att = weight*estimate) %>%
  group_by(type) %>%
  mutate(group_weight = sum(weight)) %>%
  mutate(group_att = sum(att)/group_weight) %>%
  summarise(weight = sum(weight),
            estimate = mean(estimate),
            att = mean(group_att),
            count = n(),
            overall_att = mean(full_att))

decomp_tab
```

```
## # A tibble: 3 × 6
```

##	type	weight	estimate	att	count	overall_att
##	<chr>	<dbl>	<dbl>	<dbl>	<int>	<dbl>
## 1	Earlier vs Later Treated	0.223	-4.36	-3.61	153	-5.51
## 2	Later vs Earlier Treated	0.213	-2.94	-2.62	153	-5.51
## 3	Treated vs Untreated	0.564	-6.23	-7.35	18	-5.51

Aside: Bacon Decomposition

Which shows us that the overall ATT estimated from the difference-in-differences is equal to

$$.23\hat{\beta}^{Early\ vs.\ Late} + 0.22\hat{\beta}^{Late\ vs.\ Early} + 0.56\hat{\beta}^{Treated\ vs.\ Untreated}$$

```
# Recover overall ATT as weighted combination of 2x2 comparisons:
```

```
sum(decomp_tab$weight*decomp_tab$att)
```

```
## [1] -5.507721
```

Comparing to the diff-in-diff ATT:

```
coeftable(did_1)
```

```
##              Estimate Std. Error   t value      Pr(>|t|)
## post_treat -5.507721    1.045952 -5.265751 0.0000002950744
## attr(,"type")
## [1] "Clustered (state)"
```


Aside: Bacon Decomposition

We can also plot the individual 2x2 estimates to see how much heterogeneity there is:

```
ggplot(decomp) +  
  aes(shape = factor(type)) +  
  geom_point(aes(x = weight, y = estimate), alpha = 0.4) +  
  geom_hline(yintercept = 0) +  
  labs(x = "Weight", y = "Estimate", shape = "Type") +  
  theme_minimal() +  
  theme(legend.position = "none") +  
  facet_grid(rows = vars(type), space = "fixed", scales = "fixed")
```


Aside: Bacon Decomposition

We can also plot the individual 2x2 estimates to see how much heterogeneity there is:



Aside: Bacon Decomposition

The decomposition reveals several things.

1. The most weight is given to the Treated vs. Pure Control comparison, which has the largest magnitude ATT
2. All Treated vs. Control estimates are negative
3. Lots of heterogeneity in Early vs. Late or Late vs. Early comparisons
 - Some are even positive!

Event Study

One thing that might be affecting these early/late comparisons is **dynamic treatment effects**

- We have more post-treatment observations for early adopters, fewer for late adopters
- If treatment response is dynamic, we miss that by averaging over the full post-treatment period

Estimating an event study will allow us to re-center our coefficients on **event time** and directly observe these dynamics.

Event Study

The first step is to code up an **event time** variable

- = -7 on the date 7 days prior to adoption (specific to each county)
- = -1 on the day prior to adoption
- = 0 on the date of mandate adoption
- = 1 on the first day after adoption
- = 14 on the date two weeks after adoption
- etc.

Event Time Variable

The first step is to code up an **event time** variable

```
sah_es <- group_by(sah_did, state) %>%  
  mutate(event_time = case_when(  
    # for treated states, set event_time by comparing date to  
    mandate_date (numerically)  
    !is.na(mandate_date) ~ as.numeric(date) - as.numeric(mandate_date),  
    # for controls, set event time = Inf for control units  
    TRUE ~ Inf  
  ))
```

Event Time Variable

Checking an adopting state:

```
filter(sah_es, state == "CA", event_time %in% -7:7) %>%  
  select(state, date, mandate_date, event_time)
```

```
## # A tibble: 15 × 4  
## # Groups:   state [1]  
##   state date      mandate_date event_time  
##   <chr> <date>      <date>      <dbl>  
## 1 CA    2020-03-12 2020-03-19      -7  
## 2 CA    2020-03-13 2020-03-19      -6  
## 3 CA    2020-03-14 2020-03-19      -5  
## 4 CA    2020-03-15 2020-03-19      -4  
## 5 CA    2020-03-16 2020-03-19      -3  
## 6 CA    2020-03-17 2020-03-19      -2  
## 7 CA    2020-03-18 2020-03-19      -1  
## 8 CA    2020-03-19 2020-03-19       0  
## 9 CA    2020-03-20 2020-03-19       1  
## 10 CA   2020-03-21 2020-03-19       2  
## 11 CA   2020-03-22 2020-03-19       3  
## 12 CA   2020-03-23 2020-03-19       4  
## 13 CA   2020-03-24 2020-03-19       5  
## 14 CA   2020-03-25 2020-03-19       6  
## 15 CA   2020-03-26 2020-03-19       7
```

Event Time Variable

And a control state:

```
filter(sah_es, state == "AR") %>%  
  select(state, date, mandate_date, event_time)
```

```
## # A tibble: 66 × 4  
## # Groups:   state [1]  
##   state date      mandate_date event_time  
##   <chr> <date>      <date>      <dbl>  
##  1 AR    2020-02-24 NA              Inf  
##  2 AR    2020-02-25 NA              Inf  
##  3 AR    2020-02-26 NA              Inf  
##  4 AR    2020-02-27 NA              Inf  
##  5 AR    2020-02-28 NA              Inf  
##  6 AR    2020-02-29 NA              Inf  
##  7 AR    2020-03-01 NA              Inf  
##  8 AR    2020-03-02 NA              Inf  
##  9 AR    2020-03-03 NA              Inf  
## 10 AR    2020-03-04 NA              Inf  
## # i 56 more rows
```

Event Time

Perfect! We can now see the re-centering by looking at the distribution of treatment status relative to the **date**:

```
date_freq <- sah_es %>%  
  group_by(date) %>%  
  summarise(count = sum(post_treat ==1)) %>%  
  ggplot() +  
    geom_col(aes(y = count, x = date), fill = "dodgerblue", alpha =  
0.65) +  
    geom_vline(aes(xintercept = ymd("2020-03-19")), linetype = "dashed") +  
    theme_minimal()
```




Event Time

And the distribution of treatment status relative to **event-time**:

```
eventtime_freq <- sah_es %>%  
  group_by(event_time) %>%  
  summarise(count = sum(post_treat == 1)) %>%  
  filter(event_time < Inf) %>%  
  ggplot() +  
  geom_col(aes(y = count, x = event_time), fill = "forestgreen", alpha =  
0.65) +  
  geom_vline(aes(xintercept = 0), linetype = "dashed") +  
  theme_minimal()
```

Date-time vs. Event Time

And comparing the two shows the difference under staggered adoption



- By **date**: frequency of treated states increases over time as more states adopt stay-at-home mandates
- By **event-time**: flip from no treated to fully treated at event time 0
 - Frequencies fall off for later event times as we don't have as many post-treatment event times for late adopters

Binning Endpoints

Before we calculate our dynamic treatment vector, we might want to **bin the endpoints**

- Set every event-time ≤ -24 to -24
- Set every event-time ≥ 21 to 21


This accomplishes two goals:

1. Helps with **weak identification** issues
 - i.e. only CA has event times 39-41
2. Ensures dynamic treatment effects are identified **separate from time trends** even if we omit pure control units
 - Schmidheiny and Siegloch (2019)

Binning Endpoints

Q: Why -24 and 21 for the endpoints?

A: these are the max and min event times for **all treated states**⁵



⁵ This isn't the only answer; showing robustness to alternate binning strategies would improve confidence in your estimates.

Binning Endpoints

Coding up a binned event time variable:

```
sah_es ← mutate(sah_es,  
  event_time_bin = case_when(  
    event_time ≤ -24 ~ -24, # set to -24 if ≤ -24  
    event_time ≥ 21 ~ 21, # set to 21 if ≥ 21  
    TRUE ~ event_time # if in between, keep current value  
  ))
```

Event Study

Perfect!

Now to estimate the event study, we need a **vector of indicator variables**, one for each `event_time`:

-14 Days Since SAH = 1 if event time = 14, 0 otherwise

⋮

10 Days Since SAH = 1 if event time = 0, 0 otherwise

⋮

--

We could code these by hand ...

Event Study and feols

... or we could use the **interaction operator** `i()` from **fixest** to automatically do it for us.

To do this, we need two things:

1. Our **binned event time** variable (have)
1. A **treated unit dummy** (need)

```
sah_es <- group_by(sah_es, state) %>%  
  mutate(  
    sah_state = ifelse(!is.na(mandate_date), 1, 0)  
  ) %>%  
  ungroup()
```

Event Study and feols

Now estimating the event study:

- Using `i()` to create the vector of event-time coefficients automatically
- Normalizing to the first day prior to adoption (event time -1)

```
es_reg ← feols(cadt ~ i(event_time_bin, sah_state, ref = -1) |  
               state + date, data = sah_es)
```


Event Study and feols

Looking at the output reveals all the dynamic treatment effect estimates:

```
coefstable(es_reg)
```

##		Estimate	Std. Error	t value	Pr(> t)
##	event_time_bin::-24:sah_state	5.24431750	2.2772006	2.3029669	0.0254784373305
##	event_time_bin::-23:sah_state	5.91884595	1.9222573	3.0791122	0.0033680660128
##	event_time_bin::-22:sah_state	6.03668263	1.9117363	3.1576963	0.0026955741817
##	event_time_bin::-21:sah_state	5.75074476	2.0646980	2.7852716	0.0075348287168
##	event_time_bin::-20:sah_state	5.72606565	1.8155892	3.1538333	0.0027254395246
##	event_time_bin::-19:sah_state	6.10467265	1.7414227	3.5055663	0.0009716771005
##	event_time_bin::-18:sah_state	5.97785810	1.7416719	3.4322527	0.0012101777677
##	event_time_bin::-17:sah_state	4.24352976	1.5060657	2.8176258	0.0069108599319
##	event_time_bin::-16:sah_state	4.00918806	1.3802400	2.9047035	0.0054612041340
##	event_time_bin::-15:sah_state	4.04706491	1.4052157	2.8800311	0.0058402927574
##	event_time_bin::-14:sah_state	2.73969032	1.3175311	2.0794123	0.0427277930494
##	event_time_bin::-13:sah_state	2.42408619	1.4152579	1.7128229	0.0929406324386
##	event_time_bin::-12:sah_state	2.85572278	1.5119888	1.8887196	0.0647369604521
##	event_time_bin::-11:sah_state	2.32238779	1.3237990	1.7543357	0.0855000900285
##	event_time_bin::-10:sah_state	1.02489805	1.2081339	0.8483315	0.4002964851524
##	event_time_bin::-9:sah_state	1.00625165	1.0891469	0.9238898	0.3599810611586
##	event_time_bin::-8:sah_state	1.30097049	0.9979963	1.3035825	0.1983464295810
##	event_time_bin::-7:sah_state	0.34526935	0.8934660	0.3864381	0.7008118794257
##	event_time_bin::-6:sah_state	0.42650386	0.7603438	0.5609356	0.5773467696152
##	event_time_bin::-5:sah_state	-0.02554876	0.7058742	-0.0361945	0.9712713930692
##	event_time_bin::-4:sah_state	0.16729867	0.6565861	0.2548008	0.7999220547016
##	event_time_bin::-3:sah_state	-0.16373023	0.8791889	-0.1862287	0.8530189031934
##	event_time_bin::-2:sah_state	0.78474832	0.6410018	1.2242530	0.2265941040830
##	event_time_bin::-1:sah_state	0.73371817	0.6152668	1.1844157	0.2382201770571

Event Study Plots


Unfortunately, this makes for a *super* clunky regression table.

Better: plot it!

fixest makes this super convenient with the `ipplot()` function:

```
ipplot(es_reg,  
       ylab = "Change in Average Distance Traveled (%)",  
       xlab = "Days Since Stay-at-Home in Effect",  
       main = "Event Study, Average Distance Traveled"  
)
```

iplob() Output



Event Study Plots

We can also `tidy()` up the regression output and customize the plot ourselves:

```
# get separate data frames for pre and post event times
es_pre <- tidy(es_reg) %>% filter(str_detect(term, "-"))
es_post <- tidy(es_reg) %>% filter(!str_detect(term, "-"))
# make a filler row for the reference period
es_m <- data.frame(term = "event_time_bin::-1:sah_state", estimate = 0,
std.error = 0, statistic = 0, p.value = 0)

# combine all into a single figure, add event time and confidence intervals
es_fig <- rbind(es_pre, es_m, es_post) %>%
  mutate(event_time = str_extract(term, "-?[0-9]+") %>%
as.numeric(),
         ci_l = estimate - qnorm(0.975)*std.error,
         ci_u = estimate + qnorm(0.975)*std.error)
```


Event Study Plots

If you want the fully tidyverse functionality, we can `tidy()` up the regression output and customize the plot ourselves:

```
ggplot(es_fig, aes(y = estimate, x = event_time)) +  
  # add line and points for estimates at each event-time  
  geom_line() +  
  geom_point() +  
  geom_hline(yintercept = 0, colour = "grey60", linetype = 2, alpha =  
.3) +  
  # add dashed line for SAH adoption date  
  geom_vline(xintercept = -1, linetype = "dashed", alpha = 0.7, colour =  
"navy")+  
  # add ribbon for confidence interval  
  geom_ribbon(aes(ymin = ci_l, ymax = ci_u), alpha = 0.2)+  
  theme_minimal()
```

Event Study Plots

If you want the fully tidyverse functionality, we can `tidy()` up the regression output and customize the plot ourselves:



Event Study

The event study plot makes it easy to see a few things

- **Parallel Trends:** If both groups were trending the same prior to treatment period, we'd expect event-time coefficients for $k < 0$ to be **indistinguishable from zero**
- **Dynamic Treatment Effects:** slight rebound in point estimates after a week, but statistically indistinguishable from one another here

The lack of parallel trends holding will motivate additional, forefront methods that allow us to obtain a more valid counterfactual.

Table of Contents

1. Regression Output
2. IV Regression
3. Difference-in-Differences
4. Event Study