

Lecture 9: Synthetic Control Methods, Part 2

James Sears*

AFRE 891/991 FS 25

Michigan State University

*Parts of these slides are adapted from [“Causal Inference for the Brave and True”](#) by Matheus Facure Alves.

Table of Contents

1. **Prologue**
2. **Synthetic Difference-in-Difference**
3. **Partially Pooled SCM**

Prologue

Prologue

Part 1

- Matching
- Canonical Synthetic Control

Part 2

- Synthetic Diff-in-Diff
 - Uniform Adoption
 - Staggered Adoption
- Partially Pooled Synthetic Control

Prologue

Packages we'll use today:

```
# If not installed, add in packages from GitHub not on CRAN
```

```
if (!require("augsynth"))
```

```
remotes::install_github("ebenmichael/augsynth")
```

```
if (!require("synthdid")) remotes::install_github("synth-  
inference/synthdid")
```

```
if (!require("pacman")) install.packages("pacman")
```

```
pacman::p_load(augsynth, fixest, gsynth, synthdid, tidyverse)
```

As well, let's load the event study data from [Sears et al. \(2023\)](#) we finished last lecture with (plus some covariates)

```
sah ← readRDS("data/sah_covar.rds")
```

Prologue

Last time we talked through the canonical **Synthetic Control Method (SCM)** for comparative case studies.

While it provides an avenue for overcoming violations of parallel trends, it is limited in several ways

1. Intended for a single treated aggregate unit ("comparative case studies")
2. Results primarily presented visually
3. Inference doesn't match our usual approaches

What would be *great* is if there were ways of combining the matching benefits of SCM with the identification and structure of our typical econometric approaches...

Prologue

Okay so I've already spoiled this: yes, there are methods that do this - and we're going to learn about them today.

1. Synthetic Difference-in-Differences: using a synthetic control unit as the counterfactual in a difference-in-differences
2. Partially Pooled Synthetic Control: same idea but for event study settings with dynamic treatment effects

We refer to these two approaches as doubly-robust estimators, in that you have two chances to be right:

1. In the matching design, or
2. In the regression specification

Synthetic Difference-in-Differences

Synthetic Difference-in-Differences

Synthetic Difference-in-Differences by Arkhangelsky et al. (2021) AER offers a solution to CIA violations in a typical difference-in-differences setting.

The gist:

1. Re-weight and match pre-exposure trends by constructing a synthetic counterfactual
 - Reduces reliance on parallel trends
2. Estimate a treatment effect in a standard TWFE regression

Synthetic Difference-in-Differences

More formally:

Suppose you have **balanced panel** with

- T time periods
- N units
 - First N_{co} units are never exposed to treatment (pure controls)
 - Next $N_{tr} = N - N_{co}$ units receive treatment after time T_{pre}
 - Can be block adoption or staggered adoption

You observe

- The time-varying outcome Y_{it} ,
- A binary treatment indicator W_{it} ,
- and (optionally) time-varying covariates X_{it}

i.e. same data requirements as a standard TWFE Diff-in-Diff approach

Synthetic Difference-in-Differences

The Concern: even after conditioning on observables, CIA/parallel trends doesn't hold.

The Solution: construct a synthetic control to use as the counterfactual in a Difference-in-Differences design.

Like with SCM, the goal is to find weights $\hat{\omega}^{sdid}$ so that the pre-treatment trends in the treated units' outcome align with the synthetic control's:

$$\underbrace{\sum_{i=1}^{N_{co}} \hat{\omega}^{sdid} Y_{it}}_{\text{Synthetic Control}} \approx \underbrace{\frac{1}{N_{tr}} \sum_{i=N_{co}+1}^N Y_{it}}_{\text{Treated}}$$

Synthetic Difference-in-Differences

Choosing Unit Weights: ensure average outcome for treated unit(s) is roughly parallel to the synthetic counterfactual (i.e. weighted average of control units)

Choose **unit weights** $(\hat{\omega}_0, \hat{\omega})$ that yield the arg min of

$$\underbrace{\sum_{t=1}^{T_{pre}} \left(\omega_0 + \sum_{i=1}^{N_{co}} \omega_i Y_{it} - \frac{1}{N_{tr}} \sum_{i=N_{co}+1}^N Y_{it} \right)^2}_{\text{Sum Squared Error}} + \underbrace{\zeta^2 T_{pre} \|\omega\|_2^2}_{L_2 \text{ Regularization Penalty}}$$

Synthetic Difference-in-Differences

$$\underbrace{\sum_{t=1}^{T_{pre}} \left(\omega_0 + \sum_{i=1}^{N_{co}} \omega_i Y_{it} - \frac{1}{N_{tr}} \sum_{i=N_{co}+1}^N Y_{it} \right)^2}_{\text{Sum Squared Error}} + \underbrace{\zeta^2 T_{pre} \|\omega\|_2^2}_{L_2 \text{ Regularization Penalty}}$$

- Intercept term ω_0 allows synthetic control to match on pre-trends
 - Unit FE α_i in regression step get rid of level differences
- Penalty term helps non-zero weights be more distributed across control units
 - If $\zeta = 0$, weights are identical to Abadie, Diamond, and Hainmueller (2010) SCM weights for a single treated unit
 - ζ is *complicated*, see the paper if you want more details
 - L_2 /Euclidean norm as in ridge regression (we'll chat more about this in ML lecture)

Synthetic Difference-in-Differences

Choosing Time Weights: ensure pre and post-treatment periods are balanced for control units

Choose **time weights** $(\hat{\lambda}_0, \hat{\lambda})$ that yield arg min of

$$\sum_{i=1}^{N_{co}} \left(\underbrace{\lambda_0 + \sum_{t=1}^{T_{pre}} \hat{\lambda}_t Y_{it}}_{\text{Weighted Average in Pre-Period}} - \underbrace{\frac{1}{T_{post}} \sum_{t=T_{pre}+1}^T Y_{it}}_{\text{Average in Post-Period}} \right)^2$$

- No regularization in time weights
 - Allows for correlation over time for same unit, but not across units (beyond the systematic component in a latent factor model)

Synthetic Difference-in-Differences

Once unit and time weights (and ζ) are obtained, estimate $\hat{\tau}^{sdid}$ as arg min of

$$\sum_{i=1}^N \sum_{t=1}^T (Y_{it} - \mu - \alpha_i - \beta_t - W_{it}\tau)^2 \hat{\omega}_i^{sdid} \hat{\lambda}_t^{sdid}$$

Note that without weights, this is just the TWFE Diff-in-Diff solution:

$$\sum_{i=1}^N \sum_{t=1}^T (Y_{it} - \mu - \alpha_i - \beta_t - W_{it}\tau)^2$$

Synthetic Diff-in-Diff Application

Now that we know how it's working, let's try it out by looking at the impact of California's Proposition 99, which introduced a 25 cent tax per pack of cigarettes.

```
data(california_prop99)
# colnames(california_prop99) ← "column"
# california_prop99 ←
#   separate_wider_delim(california_prop99,
#                         col = column,
#                         delim = ",",
#                         names = c("State", "Year",
# "PacksPerCapita", "treated"))
```


Synthetic Diff-in-Diff Application

We *could* evaluate the impact of Prop 99 on cigarette sales per capita using a typical TWFE Diff-in-Diff:


```
did ← feols(PacksPerCapita ~ treated | State + Year, data =  
california_prop99, cluster = ~ State)  
summary(did)
```

```
## OLS estimation, Dep. Var.: PacksPerCapita  
## Observations: 1,209  
## Fixed-effects: State: 39, Year: 31  
## Standard-errors: Clustered (State)  
##           Estimate Std. Error  t value   Pr(>|t|)  
## treated -27.3491    2.80238 -9.75925 6.6913e-12 ***  
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
## RMSE: 11.5      Adj. R2: 0.870229  
##           Within R2: 0.032671
```

Synthetic Diff-in-Diff Application

Looking at parallel trends:

```
mutate(california_prop99, state_ca = ifelse(State == "California", "California", "Pure Controls")) %>%
group_by(state_ca, Year) %>%
  summarise(PacksPerCapita = mean(PacksPerCapita)) %>%
  filter(Year < 1989) %>%
  ggplot() +
  geom_line(aes(x = Year, y = PacksPerCapita, linetype = state_ca)) +
  theme_minimal()
```



Synthetic Diff-in-Diff Application

Let's use the **synthdid** package to estimate a diff-in-diff using a synthetic California as the counterfactual.

First, we'll need to do some se the `panel.matrices()` function to set up the data

- Balanced panel (have)
- Simultaneous adoption (have mechanically)

```
synth_ca_prep <- panel.matrices(  
  panel = as.data.frame(california_prop99), # the dataframe  
  unit = "State", # unit column (name or column #)  
  time = 2, # time column (name or column #)  
  outcome = 3, # outcome var (name or column #)  
  treatment = "treated", # treatment var (name or column #)  
  treated.last = TRUE # sort treated units to be at bottom  
)
```

Synthetic Diff-in-Diff Application

Now compute the synthetic Diff-in-Diff estimate with `synthdid_estimate()`:

```
sdid ← synthdid_estimate(Y = synth_ca_prep$Y, # outcome var  
                          N0 = synth_ca_prep$N0, # number of control  
                          units (N_co)  
                          T0 = synth_ca_prep$T0 # number of pre-  
                          treatment periods (T_pre)  
                          )  
sdid
```

```
## synthdid: -15.604 +- NA. Effective N0/N0 = 16.4/38~0.4. Effective T0/T0  
= 2.8/19~0.1. N1,T1 = 1,12.
```

Which yields a much smaller treatment effect estimate than the Diff-in-Diff.

Synthetic Diff-in-Diff Application

Looking at control unit weights ω :


```
synthdid_controls(sdid)
```

```
##                estimate 1
## Nevada           0.12448923
## New Hampshire    0.10504758
## Connecticut      0.07828729
## Delaware         0.07036812
## Colorado         0.05751279
## Illinois         0.05338782
## Nebraska         0.04785319
## Montana         0.04513521
## Utah            0.04151766
## New Mexico       0.04056827
## Minnesota        0.03949464
## Wisconsin        0.03666708
## West Virginia    0.03356911
## North Carolina   0.03280518
## Idaho           0.03146821
```

Synthetic Diff-in-Diff Application

Plotting the control unit weights ω :

```
synthdid_units_plot(sdid)
```



Synthetic Diff-in-Diff Application

And the time weights λ

```
summary(sdid)$periods
```

```
##          estimate 1
## 1988          0.427
## 1986          0.366
## 1987          0.206
```

Synthetic Diff-in-Diff Application


Note that with only one treated unit we can only use the (untrustworthy) `placebo` method to get standard errors, by calling `vcov()` on our `synthdid_estimate` object:

```
se ← sqrt(vcov(sdid, method='placebo'))  
sprintf('95% CI (%1.2f, %1.2f)', sdid - 1.96 * se, sdid + 1.96 * se)  
  
## [1] "95% CI (-33.51, 2.30)"
```


Synthetic Diff-in-Diff Application

We can look at pre-treatment parallel trends by overlaying the two series:

```
synthdid_plot(sdid, overlay = 1)
```



Synthetic Diff-in-Diff Application


The built-in `synthdid` plot (a `ggplot` object) displays a lot of information by default.

- Point and line segments for simple 2x2 Diff-in-Diff comparison
- Time period weights (bottom red line)
- Customizable further with `theme()` and **other adjustments**

Synthetic Diff-in-Diff Application

The built-in `synthdid` plot (a `ggplot` object) displays a lot of information by default.

```
synthdid_plot(sdid)
```



Synthetic Diff-in-Diff Application

Comparing to canonical SCM and Diff-in-Diff reveals the differences well:


```
est_sc ← sc_estimate(Y = synth_ca_prep$Y, N0 = synth_ca_prep$N0, T0 =  
synth_ca_prep$T0)  
est_did ← did_estimate(Y = synth_ca_prep$Y, N0 = synth_ca_prep$N0, T0 =  
synth_ca_prep$T0)  
estimates ← list(est_did, est_sc, sdid)  
names(estimates) = c('Diff-in-Diff', 'Synthetic Control', 'Synthetic  
Diff-in-Diff')  
print(unlist(estimates))
```

##	Diff-in-Diff	Synthetic Control	Synthetic Diff-in-Diff
##	-27.34911	-19.61966	-15.60383

Synthetic Diff-in-Diff Application

Comparing to canonical SCM and Diff-in-Diff reveals the differences well:

```
synthdid_plot(estimates, se.method='placebo')
```



Staggered Adoption

The methods in **synthdid** easily extend to

- Controlling for covariates in the second step regression
- Cases of more than one unit adopting simultaneously.

One current limitation is lack of direct support for **staggered adoption**

Until they add native support, you can use the **Ssynthdid** package

```
# install.remotes("remotes")  
remotes::install_github("tjhon/ssynthdid")
```

Partially Pooled Synthetic Control

Partially Pooled Synthetic Control

An alternate but related estimator is the **partially pooled synthetic control method (PPSCM)** of [Ben-Michael, Feller, and Rothstein \(2021\)](#).

The gist:

- Extend canonical SCM to the many-treated unit and staggered adoption case
- Incorporate unit-level intercepts and balancing on covariates
 - Equivalent to balancing on residualized unit-level outcomes
- Estimate dynamic treatment effects
- Obtain standard errors through bootstrapping/jackknifing

Partially Pooled Synthetic Control

More formally:

Suppose you have a panel with

- T time periods
- N units
 - Some units $j = 1\dots, J$ receive treatment, potentially at different times T_i
 - Non-zero number of pure controls N_0 with $T_i = \infty$
 - Can be block adoption or staggered adoption

You observe

- The time-varying outcome Y_{it} ,
- A binary treatment indicator W_{it} ,
- and (optionally) time-varying covariates X_{it}

Partially Pooled Synthetic Control

Assumptions:

- Stable treatment and no interference across units (SUTVA)
- Prior to treatment, a unit's potential outcomes are equal to its never-treated potential outcomes
- **No Anticipation:**

$$Y_{it}(s) = Y_{it}(\infty) \text{ for } t < s, \text{ with treatment time } s$$

Partially Pooled Synthetic Control

Assumptions:

- All treated units are observed for at least several pre-periods and several post-periods
 - Needed to ensure sufficient identification in unbalanced event time
- Can express the data generating process as following
 1. Following a time-varying $AR(L)$ process:
$$Y_{it}(\infty) = \sum_{\ell=1}^L \rho_{t\ell} Y_{it-\ell}(\infty) + \epsilon_{it}, \text{ or}$$
 - Rules out correlation between treatment timing and noise terms for any period
 2. Composed of time-varying latent factors and time-invariant unit loadings: $Y_{it}(\infty) = \phi_i \cdot \mu_t + \epsilon_{it}$
 - Rules out correlation between treatment timing and noise terms *after* treatment
 3. Noise term ϵ_{it} are sub-Gaussian random variables

Partially Pooled Synthetic Control

If we want to extend canonical SCM to the **many-treated** case, we could take one of two approaches.

1. Separate SCM: estimate a separate SCM for each treated unit

- i.e. minimize pre-treatment imbalance for each treated unit separately
- Can lead to poor mean fit, biasing the ATE

2. Pooled SCM: estimate one SCM for average of treated units

- i.e. minimize the average pre-treatment imbalance across all treated units
- Can achieve strong average fit but obtain poor unit-specific treatment effects

The Solution: "pool" the two estimators, improving on each approach in isolation

Partially Pooled Synthetic Control

Choose SCM weights to minimize the weighted average of the pooled and unit-specific pre-treatment balance:


$$\min_{\Gamma \in \Delta^{scm}} \underbrace{\nu(\tilde{q}^{pool}(\Gamma))^2}_{\text{Normalized pooled imbalance}} + \underbrace{(1 - \nu)(\tilde{q}^{sep}(\Gamma))^2}_{\text{Normalized separate imbalance}} + \underbrace{\lambda ||\Gamma||_F^2}_{\text{penalize sum of squared weights}}$$

- $\hat{q}^{sep}$, $\hat{q}^{pool}$ the (normalized) root mean square of separate and pooled pre-treatment fit
- $\lambda ||\Gamma||_F^2$ a penalty term (as in SCM)
- ν the hyperparameter determining the degree of "partial pooling"
 - $\nu = 0 \Rightarrow$ Separate SCM
 - $\nu = 1 \Rightarrow$ Pooled SCM
 - $0 < \nu < 1 \Rightarrow$ Partially-Pooled SCM

Partially Pooled Synthetic Control

So, what ν should we use?


Turns out the relationship between pooled and separate imbalance is **highly convex**, with even slight interior ν offering large improvements.



Partially Pooled Synthetic Control

So, what ν should we use?

In my experience, overall ATEs are highly similar across the space of ν



PPSCM Application

So how does it actually work? Let's use the **augsynth** package and use PPSCM to revisit mobility responses to stay-at-home mandates.

Recall the setup:

- States adopted stay-at-home mandates on different days (staggered adoption)
- Observe many pre-period dates (Feb 24 to Mar 19)
- Observe many post-adoption dates (# varies by state, from 22 to 42)

PPSCM Application

One advantage to PPSCM is less data prep - we can jump directly to estimation with `augsynth()`

```
augsynth(y ~ treat | weighting covars | approx match covars | exact  
         match covars, unit, time, data, n_leads, n_lags)
```

- `formula` requires just Y_{it} and W_{it}
 - Optional weighting covariates, approximate/exact matching covariates
 - `unit/time` the names of unit and time variables
 - can be text, numeric, or dates
 - Don't need manual event time
 - `data` the dataframe
 - `n_leads/n_lags` the number of lead/lag event times to estimate
 - `n_lead` default: # post-treatment dates for last-treated unit (same as our binned event time before)
 - `n_lags` default: balance all periods

PPSCM Application

Estimating with default settings and storing the summary object

- Most everything is hanging out in the `summary()` object

```
sah_ppscm ← multisynth(cadt ~ post_treat, state, date, data = sah,  
fixedeff = TRUE)  
sum_ppscm ← summary(sah_ppscm)  
att ← sum_ppscm$att
```

PPSCM Application

Looking at the output shows

- Average ATT and standard error (across all time periods + treated units)
- Imbalances and improvement over pooled/separate SCM

```
sum_ppscm
```

```
##  
## Call:  
## multisynth(form = cadts ~ post_treat, unit = state, time = date,  
##      data = sah, fixedeff = TRUE)  
##  
## Average ATT Estimate (Std. Error): -6.585  (1.415)  
##  
## Global L2 Imbalance: 1.136  
## Scaled Global L2 Imbalance: 0.347  
## Percent improvement from uniform global weights: 65.3  
##  
## Individual L2 Imbalance: 5.543  
## Scaled Individual L2 Imbalance: 0.701  
## Percent improvement from uniform individual weights: 29.9  
##
```

PPSCM Application

The coefficient table contains two main types of data

1. "Average" period-specific ATT across all treated units (top rows, "Average" Level)
2. Unit-specific ATTs for each period

```
att_ppscm <- sum_ppscm$att  
att_ppscm
```


##	Time	Level	Estimate	Std.Error	lower_bound	upper_bound
## 1	-44	Average	-1.231576311	7.72790535	-13.06162820	11.71586842
## 2	-43	Average	1.389608170	5.82081386	-8.60717418	10.05516220
## 3	-42	Average	0.429717019	7.32960125	-12.90297642	13.30605466
## 4	-41	Average	2.236317786	5.30580965	-7.57595171	9.67595146
## 5	-40	Average	1.676580016	5.56364373	-9.06322658	8.94364911
## 6	-39	Average	1.405283071	4.41716602	-7.07210702	7.09617925
## 7	-38	Average	2.138179664	3.20732012	-4.32688271	6.99433625
## 8	-37	Average	1.286091585	2.77352618	-4.17868039	5.94519687
## 9	-36	Average	0.707602204	2.17440868	-3.66232916	4.77521443
## 10	-35	Average	0.268414075	1.70001762	-3.23701482	3.2678647

44 / 65

PPSCM Application

Plotting the Average ATT event study:


```
plot(sum_ppscm, levels = "Average")
```



PPSCM Application

Or the event study for just CA relative to its synthetic control:


```
plot(sum_ppscm, levels = "CA")
```



PPSCM Application

Or a subset (say, OR and WA):


```
plot(sum_ppscm, levels = c("OR", "WA"))
```



PPSCM Application

Or every ATT all at once:

```
plot(sum_ppscm)
```



PPSCM Application

Looking at the weights:

- $N \times N_{tr}$ matrix (51 states + DC by 43 SAH adopters)

```
sah_ppscm$weights
```

```
##           [,1]      [,2]      [,3]      [,4]      [,5]      [,6]
##           [,7]
## AK  0.0000000  0.0000000  0.000000000  0.0000000e+00  0.00000000  0.0000000e+00
##      0.00000000
## AL  0.0000000  0.0000000  0.000000000  0.0000000e+00  0.00000000  0.0000000e+00
##      0.00000000
## AR  0.0000000  0.0000000  0.000000000  5.358271e-02  0.00000000  0.0000000e+00
##      0.00000000
## AZ  0.0000000  0.0000000  0.000000000  0.0000000e+00  0.00000000  0.0000000e+00
##      0.00000000
## CA  0.0000000  0.0000000  0.000000000  0.0000000e+00  0.00000000  0.0000000e+00
##      0.00000000
## CO  0.0000000  0.0000000  0.000000000  0.0000000e+00  0.00000000  0.0000000e+00
##      0.00000000
```

PPSCM Application

What is CA (first adopter)'s synthetic control?

```
ca_wts <- sah_ppscm$weights %>%  
  as.data.frame() %>%  
  select(V1) %>%  
  rename(weight = V1)
```

```
ca_wts <- mutate(ca_wts, state = rownames(ca_wts)) %>%  
  #rename(weight = V1) %>%  
  filter(weight != 0)  
ca_wts
```

```
##      weight state  
## UT 1.000001    UT
```

```
filter(sah, state %in% c("CA", "AR", "IA", "ND", "NE", "OK", "SD", "UT",  
  "WY")) %>%  
  left_join(ca_wts, by = "state") %>%  
  mutate(state_ca = ifelse(state == "CA", "CA", "Synthetic Control"),  
    weight = ifelse(state == "CA", 1, weight)) %>%
```

PPSCM Application

Looking at ν , we see we're about halfway between separate and pooled SCM:


```
sah_ppscm$nu
```

```
## [1] 0.4902127
```

PPSCM Application

By default, the regularization penalty and number of factors for interactive fixed effect is zero. What happens if we turn them on?

```
sah_ppscm2 <- multisynth(cadt ~ post_treat, state, date, data = sah, fixedeff = TRUE,  
                        lambda = 0.5,  
                        n_factors = 2)  
sum_ppscm2 <- summary(sah_ppscm2)  
plot(sum_ppscm2, levels = "Average")
```



PPSCM Application

We can also add covariates for

- **Weighting:** covariates to weight on
- **Approx. Matching:** covariates to approximately match on *before* weighting
- **Exact Matching:** covariates to exactly match on *before* weighting


Let's weight on average pre-period values of a few covariates:

```
sah_ppscm3 <- multisynth(cadt ~ post_treat | pct_wfh + pct_pub_trans +  
  pct_retail + pct_ent_fd + poppm2019 + share_dem, state, date,  
                        data = sah, fixedeff = TRUE,  
                        lambda = 0.5,  
                        n_factors = 2)  
sum_ppscm3 <- summary(sah_ppscm3)
```

PPSCM Application

Looking at balance under this specification:

```
plot(sum_ppscm3, levels = "Average")
```




PPSCM Application

Alternatively, we can approximately match on population and democratic vote share in 2016 election using 3 nearest neighbors:

```
sah_ppscm4 ← multisynth(cadt ~ post_treat | pct_wfh + pct_pub_trans +  
  pct_retail + pct_ent_fd | poppm2019 + share_dem, state, date,  
                        data = sah, fixedeff = TRUE,  
                        lambda = 0.5,  
                        n_factors = 2,  
                        k = 3)  
sum_ppscm4 ← summary(sah_ppscm4)
```


PPSCM Application

Plotting the new Average ATTs manually with ggplot():



PPSCM Application

Comparing to the standard event study from last lecture:



Quantifying Uncertainty

Partially Pooled SCM relies upon a **wild bootstrap** for inference, but no guarantees are provided for ensuring its (asymptotic) validity.

Fortunately, [Cattaneo et al. \(2025\) REStat[†]](#) formalize methods for both **in-sample** and **out-of-sample** uncertainty in the multiple treatment unit staggered adoption SCM case that guarantee coverage probabilities non-asymptotically.

- Also: provide for estimation + inference of unit-time, unit-average, and time-average effects
- Also also: allow for a wide range of alternate weighting constraints, multiple features, and covariate adjustment
- Also also also: built an r package **scpi** to implement all this

[†] Downside: it's super new so there aren't a lot of (i.e. any) examples yet.

SCM Assumptions

We've seen that SCM approaches can yield improved parallel trends and better balance. Why is this?

A: Rather than require us to directly specify the functional form of proxies for the unobservables, they automatically test all linear combinations of the provided "input features"

Q: That's great, but what are the *right* input features?

A: Arkhangelsky and Hirshberg (2023) nicely point out that we should take advantage of all the information we have that **contains information on relevant unobservables**.

- Past outcomes are the most likely candidate for this
- If we have time-varying controls that *also* contain helpful information, then we should include those too!

SCM Assumptions

Q: So then if we have variables that we believe contain sufficient information to fully proxy for the unobservables, we're good, right?

A: Unfortunately **no**; there are still some cases where SC methods will be **biased**.

1. **Dynamic selection**
2. **Heterogeneously persistent time-varying shocks**

SCM Struggles

1. Dynamic Selection: Adoption of treatment based on past outcomes

Consider time t and a group of units that haven't adopted the treatment

- **SC Assumptions:** these units are a **good control group** for the period t
contemporaneous effect

For **dynamic effects:** need good controls in future periods $t + 1, \dots$

- **Problem:** not being treated currently $\Rightarrow$ outcomes in $t + 1, \dots$ should necessarily be systematically different

$\Rightarrow$ without additional adjustments, estimates for $t + 1, \dots$ will be **biased**

SCM Struggles

2. Heterogeneously persistent time-varying shocks

Consider a scenario where potential outcome $Y_{it}(0)$ is a function of

- Unit-level heterogeneity η_i
- Period-specific shock λ_t ,
- Some unobserved shock ϵ_{it} .

If the unobserved shock for all units arises from the **same DGP** for all units,

$$\epsilon_{it}^d = \epsilon_{it} \quad \forall i$$

then SC performs with **minimal bias** and **equal/less noise** than TWFE.

However, if there is **heterogeneity** in the persistence of the time-varying shocks, SC methods become **biased**.

SCM Struggles

2. Heterogeneously persistent time-varying shocks

If there is **heterogeneity** in the persistence of the time-varying shocks, SC methods become **biased**.

Suppose some units' potential outcomes depend on ϵ_{it}^1 and others on ϵ_{it}^2

- e.g. AR(1) vs. random walk

Then **SC fails** because there is an **added dimension of permanent heterogeneity**

- ϵ_{it}^1 , ϵ_{it}^2 matter for *both* the selection model *and* the conditional mean of counterfactual outcomes
- Linear combo of past outcomes can't distinguish units that come from ϵ_{it}^1 from those from ϵ_{it}^2
- $\Rightarrow$ **bias similar** to TWFE, **more noise**

SCM Struggles

Documentation of these struggles of SCM are **quite recent** and still limited to working papers.

The bottom line is that, while we're still figuring out the settings where SCM doesn't work, there is a large class of models in which SC methods will likely **outperform traditional TWFE estimators**.

Also, it shows an opportunity for **future research!**

- What if dynamic selection process can be modeled or well-proxied with covariates?
 - i.e. match on propensity model

Table of Contents

1. **Prologue**
2. **Synthetic Difference-in-Difference**
3. **Partially Pooled SCM**